


Things to note: these are all clean datasets. Focus of our project is not on designing a convincing model, while that is a facet of it, we want to attack the model we made and fill in our table accordingly.


MNIST - handwritten digits

- training set of 60,000 examples, and a test set of 10,000
- size-normalized and centered in a fixed-size image



CIFAR10 -

- 60000 32x32 color images in 10 classes
- 50000 training images and 10000 test images
- ID, image, label




- 10 mutually exclusive classes: airplane, automobile (but not truck or pickup truck), bird, cat, deer, dog, frog, horse, ship, and truck (but not pickup truck)

SVHN - “street view house number”

- 73257 digits for training, 26032 digits for testing, and 531131 additional, somewhat less difficult samples, to use as extra training data
- Formats:
 1. Original images with character level bounding boxes.
 2. MNIST-like 32-by-32 images centered around a single character (many of the images do contain some distractors at the sides).
- <http://ufldl.stanford.edu/housenumbers/>

Format 1:



Format 2:

